

CogniFlow - Demo Video Guide

Everything you need to record the 3-5 minute demo video for the Agentic AI Hackathon, Tech Zephyr 4.0. Written for Team BloodCoded.

You are the only speaker

This video has one narrator, and that is you. You will be doing two jobs at once - driving the terminal and explaining what appears. That is harder than it sounds, so this guide is built around making it manageable: learn one idea properly, know where each line comes from, and read the script for the rest.

There are two documents. This one teaches you the concepts and tells you *when* to say things. The second - **CogniFlow_Video_Script.pdf** - contains the actual words, split into numbered PARTs. Keep both open.

How the two documents fit together

This guide (PDF 1)	The script (PDF 2)
Explains what everything means	Gives you the exact words
Tells you which PART to read and when	Is organised into PART 0 to PART 10
Read it over the next few days	Read it while recording
Sections 1-4 teach; 5-9 prepare you	Blue panels are the only things spoken

What is in this guide

Section	What it is	When to read it
1. The one idea	What CogniFlow does, in plain English	First. Twice.
2. The five concepts	Every term you might be asked about	Over 2-3 sittings
3. Reading the screen	What each line of output means	With the app open
4. The map	Which PART to read at which moment	Before rehearsing
5. Narrating alone	Technique for doing both jobs at once	Before rehearsing
6. Setup	Exact commands, in order	Recording day
7. If it breaks	What to do live	Skim, then trust it
8. Questions	What judges may ask you	Before the finale
9. Glossary	Every term in one place	Whenever stuck
10. Learning plan	What to study, in priority order	Today

1. The one idea

If you remember nothing else, remember this page. It is the entire project, and it is what PART 1 of the script says in your own voice.

The situation

A student is learning Python. They try a recursion exercise and get it wrong. They try again and get it wrong again.

What every other AI tutor does

It gives them an **easier recursion question**. Maybe it adds a hint. Maybe it explains recursion again, more slowly.

Why that fails

Very often the student's real problem is not recursion at all. It is that they do not understand what `return` does when one function calls another. That is a **functions** problem. Recursion is just where it became visible. So the easier recursion question fails too, for exactly the same reason, and the student concludes they are 'bad at recursion' when nobody ever taught them the thing recursion is built on.

What CogniFlow does instead

It notices the pattern, works out that the real gap is **functions**, **changes its own goal** to teach functions, teaches that, and then comes back to recursion - which it never forgot was the original target.

recursion	->	functions	->	recursion
(failing)		(the real		(back to what
		gap)		they wanted)

The sentence to memorise

"Most AI tutors personalise **which question** you get. CogniFlow personalises **the route through the knowledge**."

Everything else in this guide is detail underneath that sentence. If a judge only hears one thing from you, this is the one.

2. The five concepts

These are the five things a judge might ask about. Each has: what it is, why it matters, and the one sentence to say. You do not need any mathematics.

2.1 The prerequisite graph (the core)

What it is. A map of which topic depends on which. You cannot understand recursion without functions. You cannot understand functions without variables. We wrote that map down as a structure the program can walk through.

```
variables -> functions -> recursion -> recursion_tree
variables -> loops      -> nested_loops
```

Why it matters. It lets the program ask a question no chatbot can ask: 'this student keeps failing recursion - what does recursion *depend on*, and are they weak at that instead?'

Where it appears in the video: PART 4, on the [prereq_redirect] line.

Say:

"It walks a prerequisite graph. Recursion depends on functions, so when recursion keeps failing it checks whether functions is the real problem."

2.2 Mastery, and why it is not a score out of ten

What it is. For every skill we keep a number between 0 and 1 - our best estimate of the chance the student actually knows it. Recursion at 0.35 means "we think there is about a 35% chance they have got it".

The technique is **Bayesian Knowledge Tracing**, from a 1995 paper by Corbett and Anderson. You do not need the equations. You need to know why it beats counting right answers:

- It knows people **guess**. One lucky correct answer does not prove mastery.
- It knows people **slip**. One careless mistake does not erase what they know.
- It tracks **confidence** separately. Being 80% sure after 3 questions is not the same as after 30, and the program treats those differently.

Where it appears: PART 3 and PART 6, on the [mastery] lines.

Say:

"Mastery is a Bayesian estimate, not a running total. It models guessing and slipping, and it produces a confidence number that the routing logic actually uses."

2.3 "The model proposes, the guard disposes"

What it is. Our architecture in five words, and the thing the team is proudest of. An AI model suggests what to do next - retry, move on, go back a step. It does not get the final say. Separate, ordinary, predictable code called **the guard** checks that suggestion against fixed rules and overrules it when it breaks one.

A rule the guard enforces: *you may not send a student onward if their mastery is below 0.6*. The AI might suggest it. The guard says no.

Why it matters - three reasons, worth learning properly:

Reason	What it buys us
--------	-----------------

Safety	The AI cannot wreck someone's learning path, even if it says something strange on the day.
Measurement	We can count how often the guard overrules the AI. That is a real number in our results.
Demo reliability	Because routing is predictable code, the demo takes the same path every time.

Where it appears: PART 4b - but only if a `[guard_override]` line shows up. It usually does on a live run.

Say:

"The model proposes and the guard disposes. An LLM suggests the next action, deterministic code validates it against hard rules and overrides it when it is invalid - and we log every override."

2.4 The safety invariant (our failure never costs the student)

What it is. Two kinds of bad thing can happen when a student submits code:

Their problem	Our problem
Their code has a syntax error	Our sandbox crashed
Their code crashes at runtime	The AI provider timed out
Their answer is wrong	Our database failed
Their code loops forever	The AI returned nonsense

The left column says something about the student, so it changes their mastery. The right column says something about **us**, so it must change nothing.

The clever part: this is not an `if` statement somebody has to remember to write. The two kinds are different *types* in the code, and the function that updates mastery only accepts the first kind. If an infrastructure failure ever reached it, the program would refuse outright. It is enforced by the language, not by discipline.

Where it appears: PART 5, on `[recover]` ... `mastery_untouched=True`.

Say:

"Student outcomes and system faults are disjoint types, and mastery is only reachable from the first. If our infrastructure breaks, the student does not pay for it - and that is enforced by the type system, not by a conditional someone has to notice in review."

2.5 RAG - teaching from real material

What it is. RAG stands for Retrieval-Augmented Generation. In plain terms: before the AI writes an exercise, we **search our own course notes** for the relevant section and hand it to the AI as source material.

Why it matters. Without it, the AI invents an exercise from whatever it happens to remember. With it, the exercise is grounded in the actual curriculum, and we can show which page it came from.

On screen you will see something like `05_recursion.md#5.1` The `idea` - the exact section the exercise was built from.

One detail worth knowing: we do not call RAG at every step. It runs at one point out of twelve, where it helps. If asked why, the answer is that retrieval costs time and adds nothing to arithmetic.

Where it appears: PART 3 (recursion notes) and again in PART 4, where retrieval follows the redirect to the functions chapter.

Say:

"Retrieval is filtered by skill first, then searched semantically, and every generated problem records which section it was grounded in."

3. Reading the screen

During the run, lines scroll past. Each is a real decision the program made - not a message we printed in advance. Here is how to read them.

Line you will see	What it means in plain English
[diagnose]	Working out which skill is weakest and what it depends on
[plan_action]	Choosing topic, difficulty and teaching style for the next task
[retrieve]	Searching our course notes for the right section
[generate_problem]	The AI writing the exercise, grounded in what was retrieved
[execute]	Running the student's code in a sandbox
[misconception]	Working out WHY the answer was wrong - the important one
[mastery]	Updating our estimate of what the student knows
[guard_override]	The guard rejecting the AI's suggestion - point at this
[adapt]	Deciding what to do next
[prereq_redirect]	THE MOMENT. Changing its own goal to a prerequisite
[recover]	An infrastructure failure being absorbed safely
[prereq_return]	Coming back to the original topic
[finalize]	Session finished

The three lines that matter most

- **[misconception]** - says *why* they failed, not just that they failed. Evidence that CogniFlow diagnoses causes rather than symptoms.
- **[prereq_redirect]** - the agent changing its own objective. This is the project.
- **[recover] ... mastery_untouched=True** - proof the safety invariant is real and not a claim on a slide.

4. The map - which PART to read when

This is the table you rehearse from. The words themselves are in **CogniFlow_Video_Script.pdf**.

PART	Read it when	You are doing	Length
0	Before recording anything	Setup and three test runs. Nothing spoken.	-
1	Opening shot, seeded model on screen	Pointing at functions 0.55 and recursion 0.35	30s
2	Straight after PART 1, before running	Just speaking. No screen action.	20s
3	Immediately after starting 'python run.py'	Letting the first block scroll; pointing at [diagnose], [retrieve], [mastery]	45s
4	The instant [prereq_redirect] appears	STOP SCROLLING. Hold the screen still.	50s
4b	Only if a [guard_override] line appears	Pointing at proposed= and final=	10s
5	When [execute] shows sandbox_error	Pointing at the mastery number and holding there	30s
6	When [prereq_return] appears	Scrolling to the learning path summary	25s
7	When the seven [PASS] lines are visible	Showing all seven at once	25s
8	After running 'python run.py ablation'	Showing the results table	35s
9	Final shot	Nothing. Just close.	15s
10	Only if comfortably under 5 minutes	Opening the web UI, starting a session	20s

The single most important instruction in either document

At PART 4, when [prereq_redirect] appears, **stop scrolling and leave it on screen** for the whole PART. That line is the entire project. The most common way to ruin this video is to scroll past the best moment in it.

Total timing

PARTs 1 to 9 add up to about **4 minutes 15** of speech. The rulebook allows 3 to 5 minutes, so you have roughly 45 seconds of slack for pauses and for the program thinking. That is comfortable - but it is not enough slack to also include PART 10 unless you are running fast. **Going over five minutes is worse than leaving PART 10 out.**

5. Narrating alone

You are doing two jobs at once. These techniques make that manageable. Choose your approach during rehearsal, not on the day.

Choose one of these two approaches

Approach	How it works	Trade-off
A - Narrate live	Start the command and talk while output scrolls. No expressive pauses that keep you from the output and there are no corrections.	Less expressive, pauses that keep you from the output and there are no corrections.
B - Run first, then narrate the whole run finish. Then scroll back through the screen and narrate the pace it ran by. The output is complete, and you control the pace it ran by.	Start the command and talk while output scrolls. No expressive pauses that keep you from the output and there are no corrections.	Less expressive, pauses that keep you from the output and there are no corrections.

Approach B is not cheating

The rulebook forbids fabricating or manipulating **results**. Scrolling back through the genuine output of a genuine run and explaining it is not that - it is how every code walkthrough in the world is done. What would be wrong is inventing output, or claiming something the screen does not show. **If you are recording alone and nervous, take approach B.** A calm, clear video beats a rushed, authentic-but-flustered one.

Recording in segments

You do not have to do this in one take. Record each PART separately, stop, breathe, then start the next. Join them afterwards in any free editor. This is the single biggest stress reducer available to you, and no one can tell from the finished video.

Six delivery habits

- **Point with the cursor** at the line you are describing. It tells the viewer where to look and it keeps you anchored to what is really on screen.
- **Pause after the important sentences.** Silence feels much longer to you than to a viewer. Two seconds after 'It does not move' in PART 5 is right.
- **Slow down at PART 1 and PART 4.** Those two carry the whole idea. Everything else can be brisker.
- **Do not read in a monotone.** The script is written the way people speak. Say it the way you would explain it to a friend, not the way you would read a notice.
- **If you stumble, stop and redo that PART.** Do not push through - a clean retake costs 30 seconds.
- **Never say a thing the screen is not showing.** If in doubt, describe exactly what is visible. That is always safe and always true.

6. Setup - recording day

This is PART 0 of the script, repeated here so it is in both documents.

Step 1 - Prove it works before recording anything

```
cd D:\Downloads\BloodCoded
python run.py check      # one real call per AI provider
python run.py live       # the full demo against real models
```

You want **ALL CHECKS PASSED** and the line `Learning path : recursion -> functions -> recursion`.

Step 2 - Run it three times

Run `python run.py live` three times and confirm the learning path is **identical** each time. It will be - routing is deterministic by design. If it varies, stop and tell Sohan before recording. A demo you cannot reproduce is a demo that will fail in front of judges.

Step 3 - Time it

While you are doing those three runs, **time how long the run takes**. That number decides whether you take approach A or B in section 5, and it is the one thing this guide cannot tell you in advance because it depends on how fast the AI providers respond on the day.

Step 4 - Screen and audio

- Terminal about 110 characters wide, font large enough to read on a phone.
- Close everything else: notifications, browser tabs, chat apps.
- Dark terminal theme reads better on video than light.
- Record ten seconds and play it back before a real take.
- Keep the script PDF on a second screen or printed - not on the screen you record.

7. If something breaks while recording

The rule

Say what happened, plainly, and keep going. Do not pretend it did not happen and do not talk over it. A recovered failure demonstrates the recovery path we are claiming - genuinely better evidence than a run where nothing went wrong. And if a take is spoiled, just record that PART again.

What you see	What it means	What to say
degraded=True on a line	An AI provider failed, so a built-in template was used	The template was used
provider=google:...	The first provider hit its limit; it failed to schedule a provider	"That is a live API call - this is not a recording."
A long pause	Waiting for a real API response	"That is a live API call - this is not a recording."
An error you do not recognise	Unknown	"Something went wrong there - let me re-run it." Then run 'python run.py verify'
The path came out different	Should not happen - routing is correct	Stop. Tell Sohan before continuing.

Known quirk worth understanding

Our first AI provider (Groq) has a free limit of about 8,000 tokens per minute, and one full run uses roughly 35,000. So it **will** hit the limit partway through, and it **will** automatically switch to the second provider (Google). You will see `provider=google:gemini-3.5-flash-lite` appear. That is not a fault - it is the failover working, and it is worth pointing out if it happens.

8. Questions a judge may ask

Short honest answers beat long ones. **"I do not know, but I can show you where it is in the code" is a perfectly good answer** and far better than guessing.

Question	Your answer
Is the redirect hardcoded?	No. It is graph traversal over live mastery estimates. There is a test that asserts the state transition
Why not just use a big LLM?	An LLM has no calibrated model of the student, so it will happily advance someone who is not ready
Why not just rules, then?	Rules cannot author a novel exercise grounded in a specific misconception and a specific page of
Is the sandbox real?	Yes - a separate process with a timeout and a minimal environment. Student code cannot read or
Is mastery just a counter?	No, it is Bayesian Knowledge Tracing. It models slipping and guessing and yields a confidence score
Did you train a model?	We built parameter fitting for the mastery model. It produced a negative result on held-out data, so
What happens if the AI is down?	It fails over to a second provider. If all fail, it uses a deterministic template and marks the event de
Did you write this yourselves?	Yes. Two contributors, both human, and the contributor list is checked automatically before every
What does not work yet?	There is no login, no dashboard and no diagnostic quiz - a new student starts from a seeded prof
What breaks it?	Seven and a half percent of students with no gap still get an unnecessary detour. It was twenty-fi

9. Glossary

Term	Plain meaning
Agent	A program that decides its own next step instead of following a fixed script.
LangGraph	The library we use to define those steps and the paths between them.
Node	One step in that structure - 'diagnose', 'grade', 'adapt'.
State	Everything the program currently knows about this session.
Checkpoint	A saved copy of that state, written to disk, so it can resume later.
Interrupt	The program genuinely stopping to wait for the student - not a loop that spins.
Mastery	Our 0-to-1 estimate of whether the student knows a skill.
Confidence	How much evidence that estimate rests on.
BKT	Bayesian Knowledge Tracing - the method behind those two numbers.
Prerequisite	A skill you must know before another one makes sense.
The guard	Predictable code that checks, and can overrule, the AI's suggestion.
RAG	Searching our own notes and giving them to the AI as source material.
Sandbox	The isolated place student code runs, so it cannot harm anything.
SystemFault	A failure that is ours, not the student's. Never changes mastery.
Ablation	An experiment where you remove one part to prove it was doing something.
Deterministic	Same input, same output, every time. No randomness.
Degraded	A step that fell back to a built-in template because a provider failed.

10. What to study, in priority order

You do not need all of this. If you only do the first three rows, you can still record a good video.

Priority	What	Time	How
Essential	Section 1, until you can say the one idea without reading it	30 min	Read it, then explain it out loud to someone
Essential	Watch a run three times, following the lines in section 3 'python run.py live' with this guide open beside you	45 min	
Essential	Read the whole script PDF once, out loud	30 min	Out loud matters - silent reading hides the awkward sentences
Essential	Rehearse PARTs 1 to 9 with the screen, twice	1.5 hours	Time it. Aim for 4:15 of speech.
Strongly	Sections 2.3 and 2.4 - the guard and the safety invariant	15 min	These two are what a judge is most likely to probe
Strongly	The Q&A table in section 8	30 min	Say the answers aloud; do not just read them
Strongly	Section 5 - decide approach A or B	15 min	Decide in rehearsal, not on the day
If time	Section 2.2 - what BKT actually does	45 min	Optional. The one-sentence version is enough.
If time	Open app/mastery/policy.py and read the rules	15 min	Ordinary Python with comments. Less scary than it sounds.

A final thought

The strongest thing you can do in a technical Q&A; is be straight about the limits. We know what CogniFlow does not do yet - no login, no dashboard, no diagnostic quiz - and saying so plainly makes every other claim more believable, not less. Confidence is knowing where the edges are, not pretending there are none.

Good luck. The engine works. All you have to do is show it honestly.